


# AlphaGo Zero

Silver D., Schrittwieser J., Simonyan I., et al. 2017



# Overview

- ❑ Cool Findings
- ❑ Neural Network Architecture
- ❑ Search Algorithm
- ❑ Training Pipeline
- ❑ Analysis

# Overview

❑ Cool Findings

❑ Neural Network Architecture

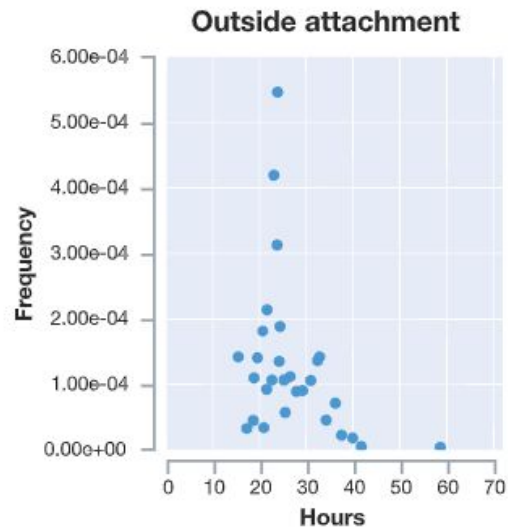
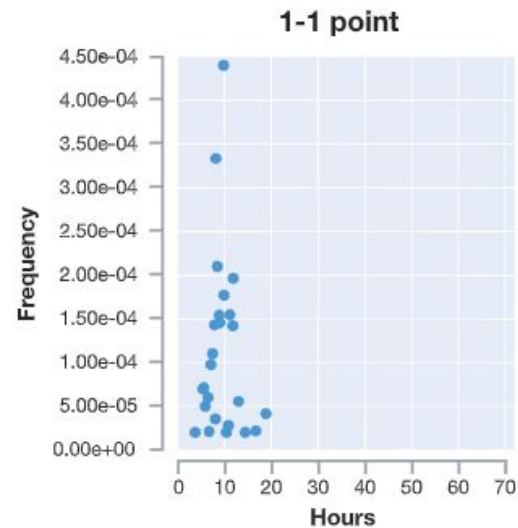
❑ Search Algorithm

❑ Training Pipeline

❑ Analysis

# Cool Findings



- Discovery of common human moves, then later discarding them for novel moves





# Cool Findings

- Discovery of common human moves, then later discarding them for novel moves

3-3 invasion




3-3 point knight's move



# Cool Findings

- ❏ Number of moves considered, human vs AlphaZero vs Stockfish



# Overview

- ❑ Cool Findings
- ❑ Neural Network Architecture
- ❑ Search Algorithm
- ❑ Training Pipeline
- ❑ Analysis

# Neural Network Architecture: Inputs

$$s_t = [X_t, Y_t, X_{t-1}, Y_{t-1}, \dots, X_{t-7}, Y_{t-7}, C]$$

$[X_t, \dots, X_{t-7}]$  19 X 19 binary feature plane of current players stones over last 8 moves.

$[Y_t, \dots, Y_{t-7}]$  19 x 19 binary feature plane of opponents stones over last 8 moves.

$[C]$  Handicap points to indicate who's turn it is.



# Neural Network Architecture: Outputs

$p_t$

**Policy vector**

Probability distribution over all moves including pass.

- Probability corresponds to how good the move is
- $19 \times 19 + 1 = 362$

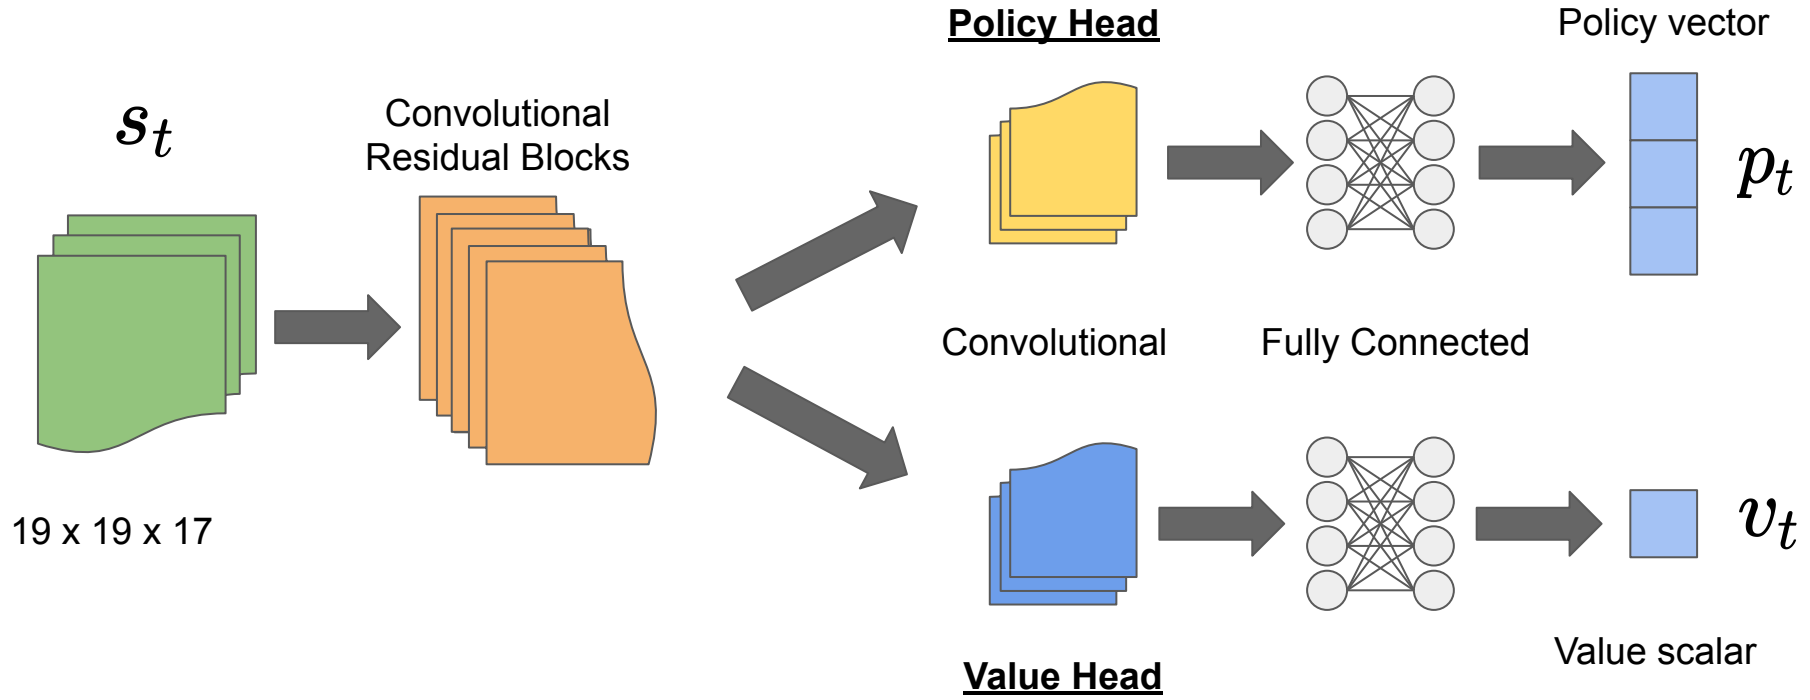
$v_t$

**Value scalar**

Scalar indicating the likelihood of winning from current state.

- $[-1, 1]$

# Neural Network Architecture



# Neural Network Architecture

$$f_{\theta}(s_t) = (p_t, v_t)$$

$$l = (z - v)^2 - \pi^T \log(p) + c ||\theta||^2$$

# Overview

- ❑ Cool Findings
- ❑ Neural Network Architecture
- ❑ Search Algorithm
- ❑ Training Pipeline
- ❑ Evaluation

# Search Algorithm

## **Vanilla Monte Carlo Tree Search (MCTS)**

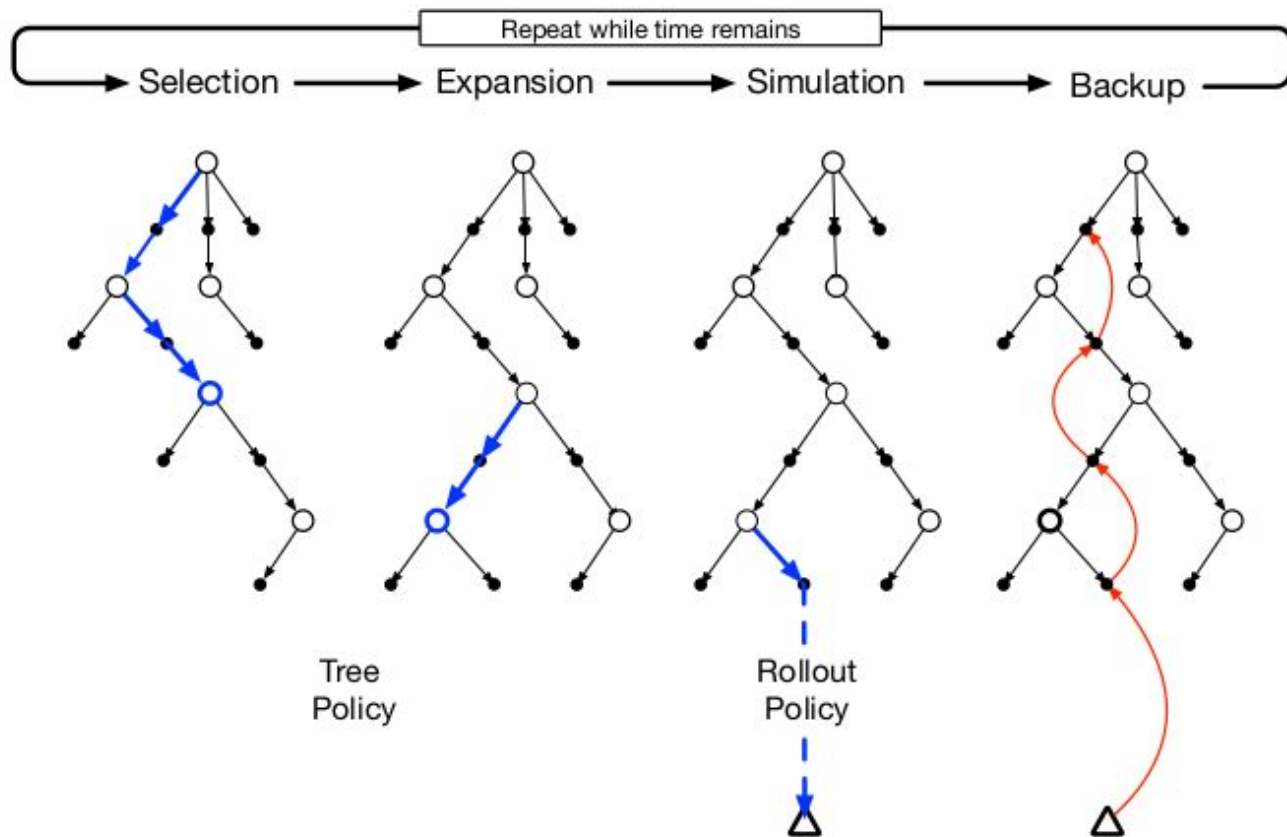
- ❑ Coulom, R. 2006, Kocsis, L., Szepesvari, C. 2006.
- ❑ Utilizes multiple simulated trajectories from current state to pick best action
- ❑ Focus successive simulations to extend trajectories of high value
- ❑ Effective when environment model is simple enough for fast simulation

# Search Algorithm

## Vanilla Monte Carlo Tree Search (MCTS)

1. **Selection** - start at root and choose the best path
2. **Expansion** - once at leaf node, expand one or more child nodes
3. **Simulation** - simulate trajectory until terminal state
4. **Backup** - use terminal value to update tree

# Search Algorithm



# Search Algorithm

## AlphaGo Zero (MCTS)

- ❑ Utilizes the neural network to guide search

$$f_{\theta}(s_t) = (p_t, v_t)$$

- ❑ Reduces the breath of the search with the policy vector

$$p_t$$

- ❑ Reduces the depth of the search the value scalar

$$v_t$$



# Search Algorithm

## AlphaGo Zero (MCTS)

- ❑ Each edge  $(s, a)$  stores statistics:  $\{N(s, a), W(s, a), Q(s, a), P(s, a)\}$ 
  - ❑ Visit count  $N(s, a)$
  - ❑ Total action value  $W(s, a)$
  - ❑ Mean action value  $Q(s, a)$
  - ❑ Prior probability  $P(s, a)$

# Search Algorithm

## AlphaGo Zero (MCTS)

1. **Selection** - Choose edges according to edge statistics

$$a_t = \underset{a}{\operatorname{argmax}} (Q(s_t, a) + U(s_t, a))$$

- ❑ Q is the **exploitative** term
- ❑ U is the **exploration** term

$$U = c_{punct} P(s, a) \frac{\sqrt{\sum_b N(s, b)}}{1 + N(s, a)}$$

# Search Algorithm

## AlphaGo Zero (MCTS)

### 1. Selection

$$U = c_{punct} P(s, a) \frac{\sqrt{\sum_b N(s, b)}}{1 + N(s, a)}$$

- ❑ Exploration constant  $c_{punct}$
- ❑ Prior probability  $P(s, a)$
- ❑ Parent visit count  $\sum_b N(s, b)$
- ❑ Edge visit count  $N(s, a)$

# Search Algorithm

## AlphaGo Zero (MCTS)

2. **Expand and evaluate** - once a leaf node is reached, expand node with NN

$$f_{\theta}(s_t) = (p_t, v_t)$$

- ❑ Each edge is initialized:

$$\{N(s, a) = 0, W(s, a) = 0, Q(s, a) = 0, P(s, a) = p_t\}$$

- ❑ The value is used for back up:  $v_t$

# Search Algorithm

## AlphaGo Zero (MCTS)

3. **Backup** - edge statistics are update in backward pass to the root

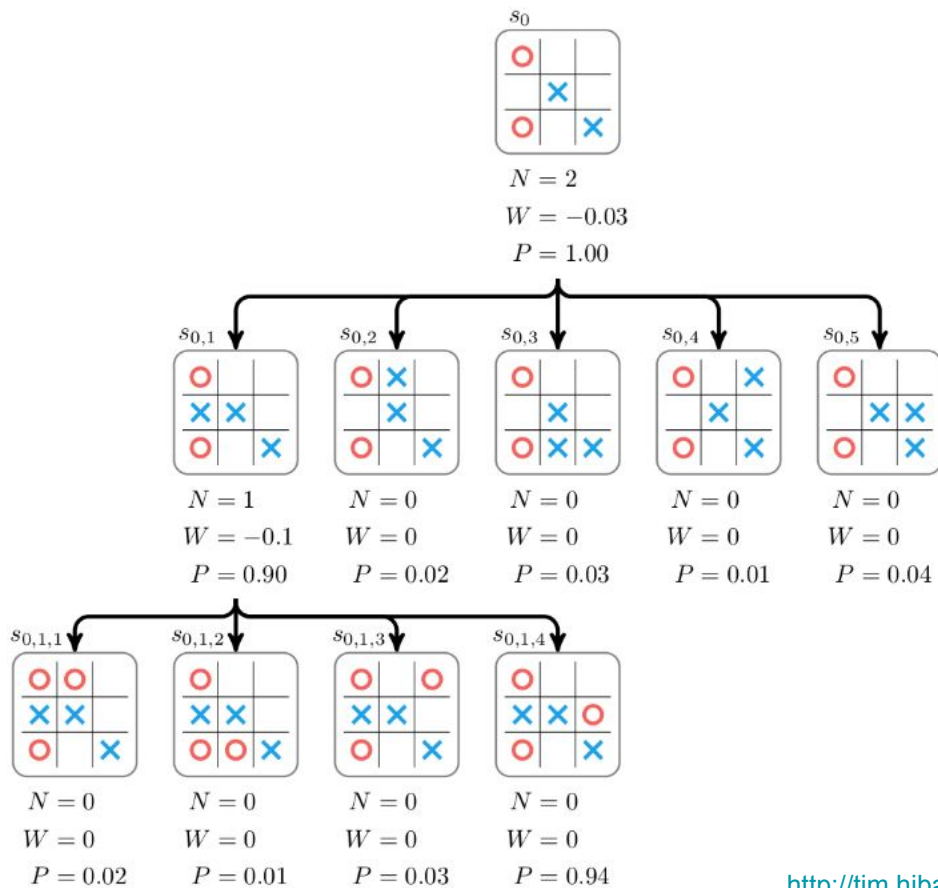
- ❑ Node visit counts are incremented:

$$N(s_t, a_t) = N(s_t, a_t) + 1$$

- ❑ Action value is updated to the mean value:

$$W(s_t, a_t) = W(s_t, a_t) + v, \quad Q(s_t, a_t) = \frac{W(s_t, a_t)}{N(s_t, a_t)}$$

# Search Algorithm



# Search Algorithm

## AlphaGo Zero (MCTS)

4. **Play** - choose action  $a$  to play in root position  $s_0$

$$\pi(a \mid s_0) = \frac{N(s_0, a)^{1/\tau}}{\sum_b N(s_0, b)^{1/\tau}}$$

- ❑ Temperature parameter, controls exploration:  $\tau$
- ❑ Edge visit count:  $N(s_0, a)$
- ❑ Root visit count:  $\sum_b N(s_0, b)$

# Search Algorithm

## AlphaGo Zero (MCTS)

4. **Play** - save MCTS policy, state, and value placeholder for NN training

$$(s_t, \pi_t, z_t)$$

- At end of game, the value placeholder is updated with final reward:

$$z_t = r_T = \{-1, +1\}$$




# Overview

- ❑ Cool Findings
- ❑ Neural Network Architecture
- ❑ Search Algorithm
- ❑ Training Pipeline
- ❑ Analysis

# Training Pipeline

We have a NN and we can use it to play games. How do we make it better?



# Training Pipeline: Review: Iterative Policy Evaluation

## 1. Policy Iteration:

a. **Policy Evaluation:** Estimate value following policy



$$V = V_{\pi}$$

b. **Policy Improvement:** Generate new policy


$$\pi' \geq \pi$$

c. **Repeat:** Continue until convergence

$$V^*, \pi^*$$



# Training Pipeline : Review: Iterative Policy Evaluation



- Undiscounted episodic MDP ( $\gamma = 1$ )
- Nonterminal states  $1, \dots, 14$
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is  $-1$  until the terminal state is reached
- Agent follows uniform random policy

$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$$

\*These slides are horribly ripped off from David Silver's excellent "Introduction To Reinforcement Learning" course <https://www.youtube.com/playlist?list=PLqYmG7hTraZDM-OYHWgPebj2MfCFzFObQ>

# Training Pipeline : Review: Bellman Equation

The value function can be decomposed into two parts:

- immediate reward  $R_{t+1}$
- discounted value of successor state  $\gamma v(S_{t+1})$

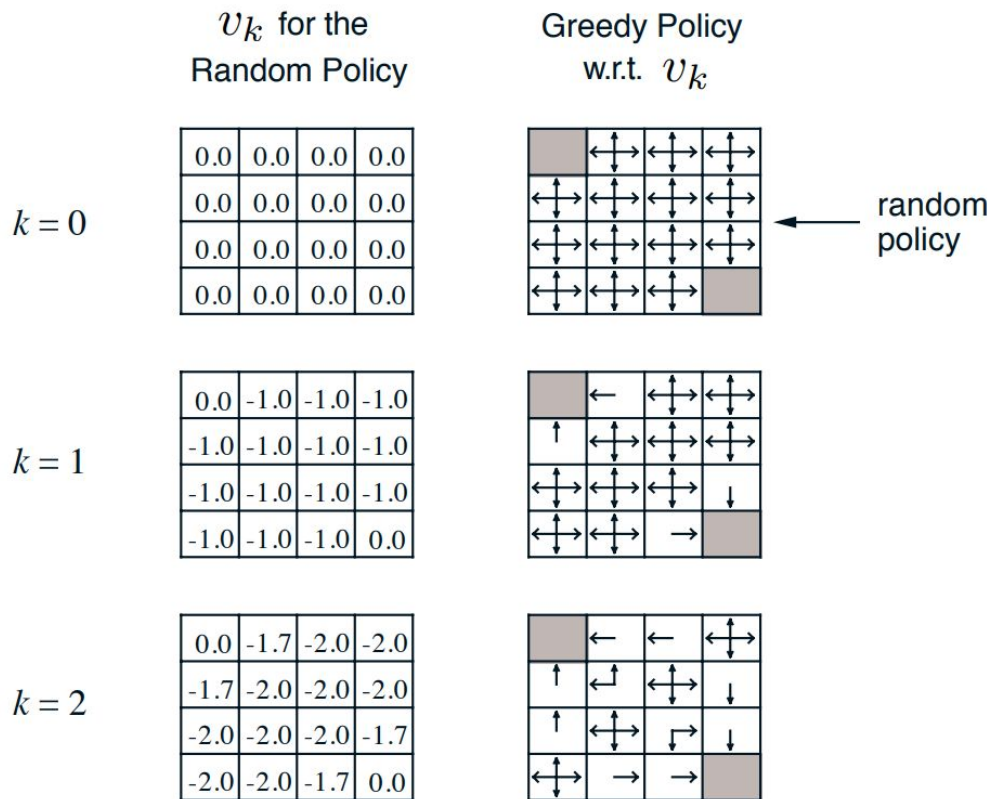
## Definition

The *return*  $G_t$  is the total discounted reward from time-step  $t$ .


$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$\begin{aligned} v(s) &= \mathbb{E}[G_t \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma (R_{t+2} + \gamma R_{t+3} + \dots) \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s] \end{aligned}$$

# Training Pipeline : Review: Iterative Policy Evaluation



# Training Pipeline : Review: Iterative Policy Evaluation



# Training Pipeline

Training is comprised of three components:

1. **Optimization:** Neural Network optimized on recent self-play data
2. **Evaluator:** AlphaGo players are evaluated against current best player
3. **Self-play:** Best current player is used to generate self-play data



# Training Pipeline

## 1. Policy Iteration:

### a. **Policy Evaluation:** $(s_t, \pi_t, z_t)$

- MCTS takes a sampling based approach (using priors to guide search) to finding the true value of your sub-MDP then uses that information to make an improved move

### b. **Policy Improvement:**

- “Policy improvement starts with a neural network policy, executes an MCTS based on that policy’s recommendations, and then projects the (much stronger) search policy back into the function space of the neural network.”

### c. **Repeat:**

- Self play, train  $f_{\theta}(s_t) = (p_t, v_t)$

$$l = (z - v)^2 - \pi^T \log(p) + c ||\theta||^2$$

# Training Pipeline

## 2. Evaluator:



# Training Pipeline

## 2. Evaluator:

- ❑ Each NN checkpoint is evaluated against current best player:  $f_{\theta_i}$  vs  $\alpha_{\theta_*}$
- ❑ MCTS is used to pick moves with:  $\tau \rightarrow 0$
- ❑ If new player wins  $> 55\%$  of games then:  $\alpha_{\theta_*} \leftarrow f_{\theta_i}$

# Training Pipeline

## 3. Self-play:



# Training Pipeline

## 3. Self-play:

- ❑ Current best player is used to generate data

- ❑ For the first 30 moves:  $\tau \rightarrow 1$

- ❑ After that:  $\tau \rightarrow 0$

$$\pi(a \mid s_0) = \frac{N(s_0, a)^{1/\tau}}{\sum_b N(s_0, b)^{1/\tau}}$$

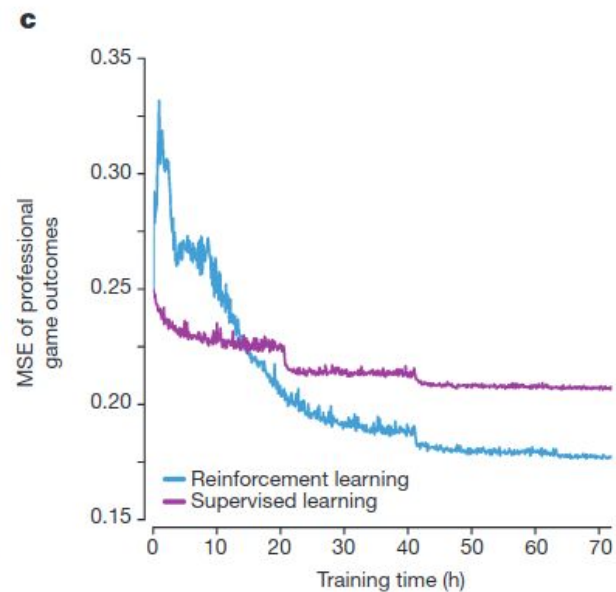
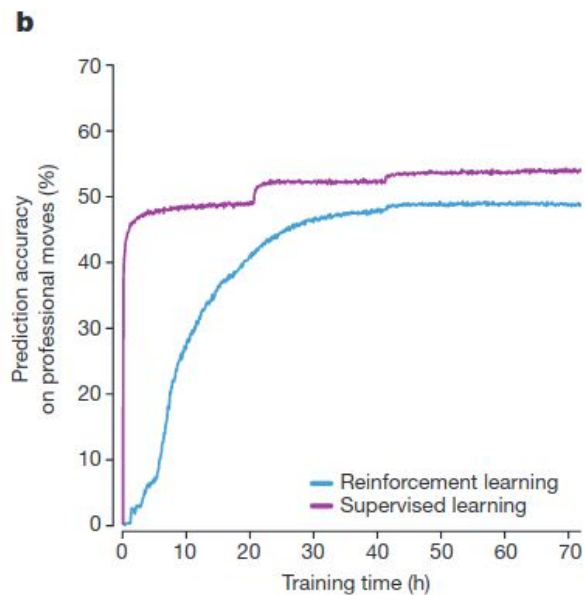

- ❑ Dirichlet noise is added to prior probabilities in MCTS:

$$P(s, a) = (1 - \varepsilon)p_a + \varepsilon\eta \quad \varepsilon = 0.25, \eta \approx \text{Dir}(0.03)$$

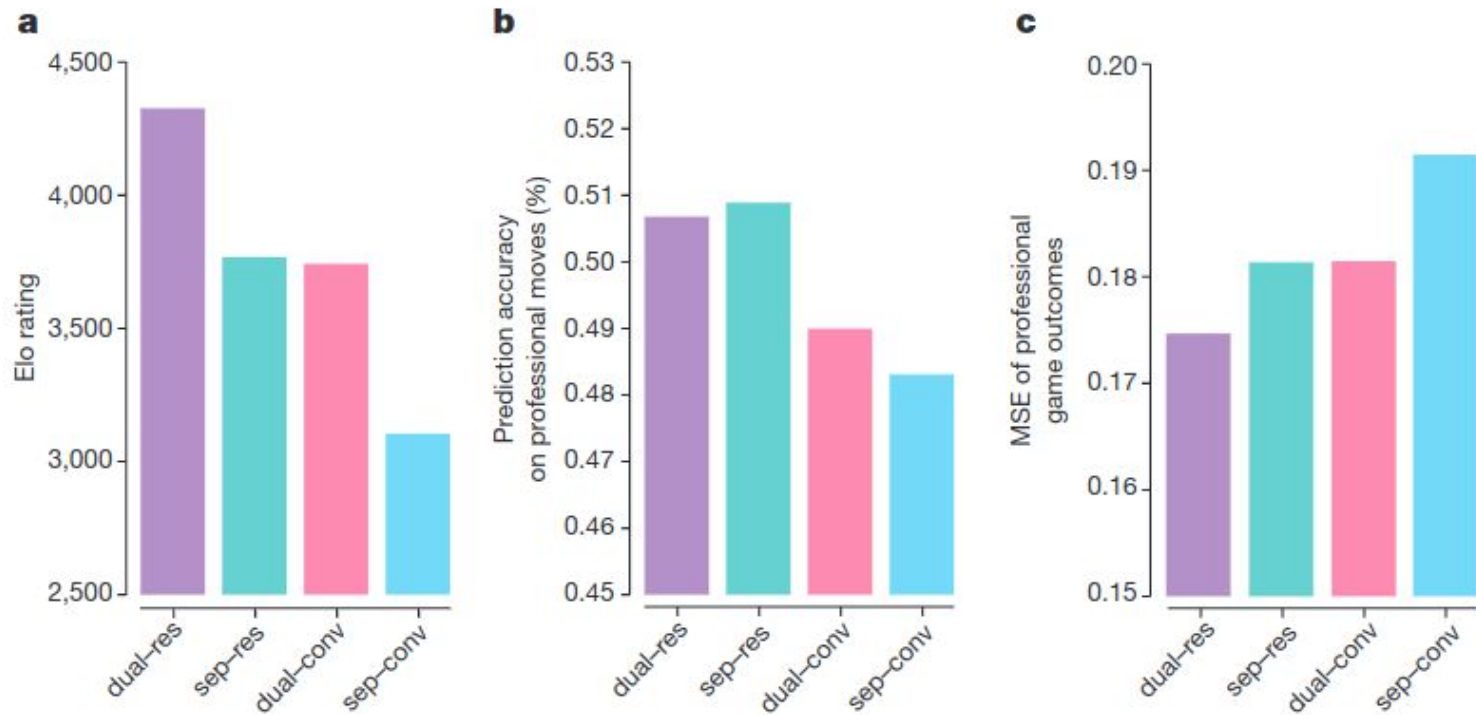
# Overview

- ❑ Cool Findings
- ❑ Neural Network Architecture
- ❑ Search Algorithm
- ❑ Training Pipeline
- ❑ Analysis

## Supervised Learning vs Self-play



## Architecture Comparison





# Analysis

## Performance Comparison

